

Variational Auto-Encoder and generative Zoo

Machine Learning and Data Mining

Maxim Borisyak

National Research University Higher School of Economics (HSE)

November 29, 2017

Generative models

Generative models

- Informally, given samples we wish to learn generative procedure.
- Formally, given samples of a random variable X , we wish to find X' , so that:

$$P(X) \approx P(X')$$

In the previous episodes

RBM:

- Maximum Likelihood fit through energy function;
- Gibbs sampling;

GAN:

- transformation from a tractable random variable to a target $X = f(Z)$;
- minimizes Jensen-Shannon distance which estimated via classifier;
- direct sampling.

Variational Auto-Encoder

Latent variables revisited

Before generating a sample, model should first decide what it should generate:

- which digit to generate: 0, 1, ..., 9
- width of stokes;
- 'speed';
- etc.

Such decision can be represented as random variables, often called **latent variables**.

Latent variables revisited

Like most of the generative models, VAE searches random variable as a function of *latent variables* Z :

- easy to sample;
- tractable distribution.

Most common choice $Z \sim \mathcal{N}^n(0, 1)$. Unlike GAN, transformation $Z \rightarrow X$ is **not deterministic**:

$$\begin{aligned}P(X) &= \int P(X | z)P(z)dz; \\P(X|z) &= \mathcal{N}(X | f(z), \sigma^2 I).\end{aligned}$$

$$P(X) = \int P(X | z)P(z)dz = \mathbb{E}_z P(X | z)$$

- sampling from $P(X)$ might be computationally expensive;
- most of z produce $P(X | z) \approx 0$.

Variational bound

$$P(X) = \int P(X | z)P(z)dz = \mathbb{E}_Z P(X | Z)$$

In order to make sampling tractable, $P(Z)$ can be replaced by some $Q(Z | X)$.

$$\log P(X) = \text{KL} (Q(Z | X) \| P(Z | X)) + \mathcal{L}(X) \geq \mathcal{L}(X)$$

$$\begin{aligned} \mathcal{L}(X) &= \mathbb{E}_{Z \sim Q} [-\log Q(Z | X) + \log P(X, Z)] = \\ &\quad - \text{KL} (Q(Z | X) \| P(Z)) + \mathbb{E}_{z \sim Q(Z|X)} [\log P(X | Z)] \end{aligned}$$

VAE objective

$$\log P(X) - \text{KL}(Q(Z | X) \| P(Z | X)) =$$
$$\mathbb{E}_{z \sim Q(Z|X)} \log P(X | Z) - \text{KL}(Q(Z | X) \| P(Z))$$

where:

- $\text{KL}(Q(Z | X) \| P(Z | X))$ - error term:
 - recognition model penalty;
- $\mathbb{E}_{z \sim Q(Z|X)} [\log P(X | Z)]$ - reconstruction error:
 - can be estimated like in an ordinary AE;
- $\text{KL}(Q(Z | X) \| P(Z))$ - something similar to regularization;
 - can be computed analytically if $P(Z)$ is well defined.

Reconstruction error

$$\text{RE} = \mathbb{E}_{z \sim Q(Z|X)} [\log P(X | Z)]$$

- for Gaussian posterior i.e. $P(X | Z) = \mathcal{N}(X | f(z), \sigma^2 I)$:

$$\text{RE} = \frac{1}{2} E_{Z \sim Q(Z|X)} [f(Z) - X]^2 + \text{const}$$

- for Bernoulli posterior (e.g. for discrete output)

$$P(X = 1 | Z) = f(z):$$

$$\text{RE} = E_{Z \sim Q(Z|X)} [X \log f(Z) + (1 - X) \log(1 - f(Z))]$$

The other term

$$\text{KL}(Q(Z | X) \| P(Z))$$

Consider:

- $Q(Z | X) = \mathcal{N}(Z | \mu(X), \Sigma(X));$
- $P(Z) = \mathcal{N}(0, I)$

$$\begin{aligned} \text{KL}(\mathcal{N}(X | f(Z), \Sigma(Z)) \| \mathcal{N}(X | f(Z), \Sigma(Z))) = \\ \frac{1}{2} (\text{tr}(\Sigma(X)) + \|\mu(X)\|^2 - k - \log \det \Sigma(X)) = \\ \frac{1}{2} \left(\|\mu(X)\|^2 + \sum_i \Sigma_{ii}(X) - \log \Sigma_{ii}(X) \right) - \frac{k}{2} \end{aligned}$$

Training time

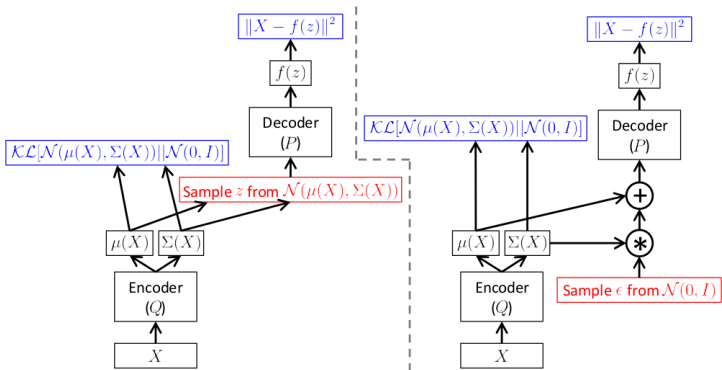


Figure 4: A training-time variational autoencoder implemented as a feed-forward neural network, where $P(X|z)$ is Gaussian. Left is without the “reparameterization trick”, and right is with it. Red shows sampling operations that are non-differentiable. Blue shows loss layers. The feedforward behavior of these networks is identical, but backpropagation can be applied only to the right network.

Testing time




Figure 5: The testing-time variational “autoencoder,” which allows us to generate new samples. The “encoder” pathway is simply discarded.

Conditional VAE

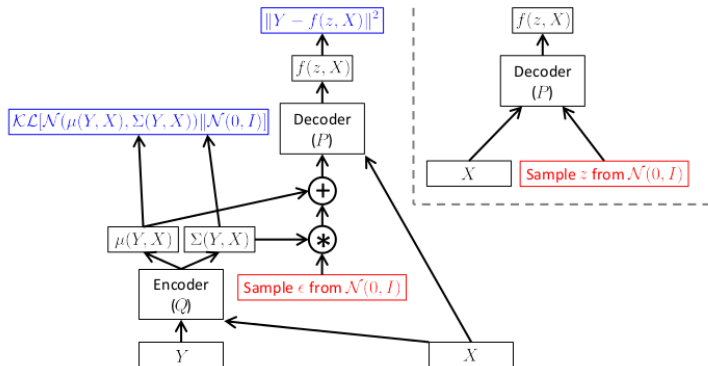


Figure 6: Left: a training-time conditional variational autoencoder implemented as a feedforward neural network, following the same notation as Figure 4. Right: the same model at test time, when we want to sample from $P(Y|X)$.

Discussion

VAE:

- simple to implement;
- stable training;
- fast training;
- form of inference and generative network impose serious restrictions on possible models;
- might be tricky to modify.

GAN:

- might be tricky to implement;
- training requires tricks (or EB/W-GAN);
- training is slow;
- often, produces state-of-the-art models;
- flexible architecture.

BiGAN

The main difference between GAN and VAE is absence of an inference model in GAN ($Q(Z|X)$ in VAE). BiGAN adds this inference by introducing an inference networks:

$$X' = G(Z);$$

$$Z' = E(X);$$

and solving min-max problem for:

$$\mathcal{L}(D, E, G) = \mathbb{E}_x \log D(x, E(x)) + \mathbb{E}_z \log(1 - D(G(z), z))$$




Figure 1: The structure of Bidirectional Generative Adversarial Networks (BiGAN).

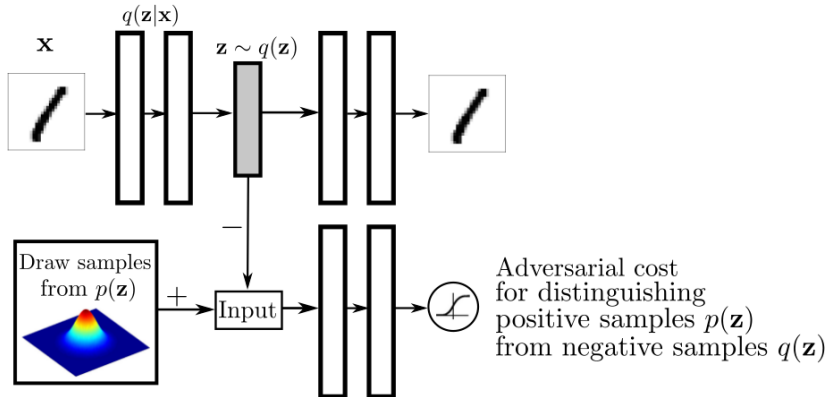
- BiGAN can be used as dimensionality reduction technique;
- unlike AE, importance of a feature produced not by Euclidean distance, but by a discriminator.

Adversarial AutoEncoder

Adversarial AutoEncoder

- replaces VAE regularization term with GAN objective;
- also makes both inference and generative networks deterministic:
 - removes possible error source;
- essentially, AutoEncoder with regularization term that drives code space to a predefined distribution.

Adversarial AutoEncoder



Adversarial AutoEncoder

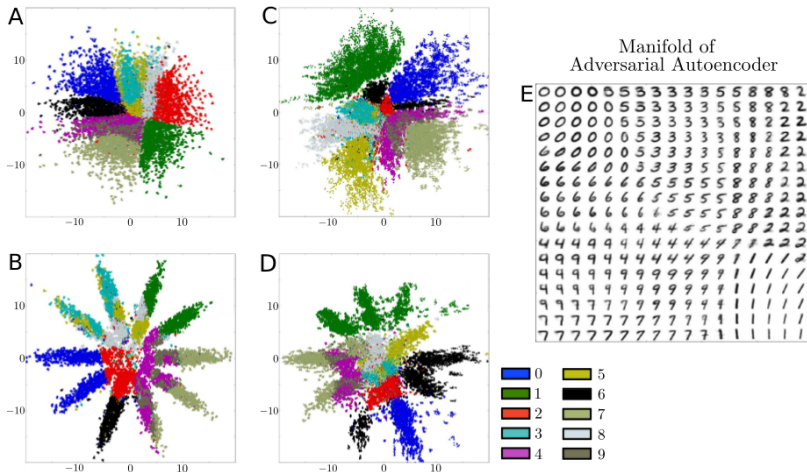


Figure 2: Comparison of adversarial and variational autoencoder on MNIST. The hidden code z of the *hold-out* images for an adversarial autoencoder fit to (a) a 2-D Gaussian and (b) a mixture of 10 2-D Gaussians. Each color represents the associated label. Same for variational autoencoder with (c) a 2-D gaussian and (d) a mixture of 10 2-D Gaussians. (e) Images generated by uniformly sampling the Gaussian percentiles along each hidden code dimension z in the 2-D Gaussian adversarial autoencoder.

Adversarial AutoEncoder

Codes can be associated with known high level features, e.g. face expression.




Figure 6: Disentangling the label information from the hidden code by providing the one-hot vector to the generative model. The hidden code in this case learns to represent the style of the image.

Adversarial AutoEncoder

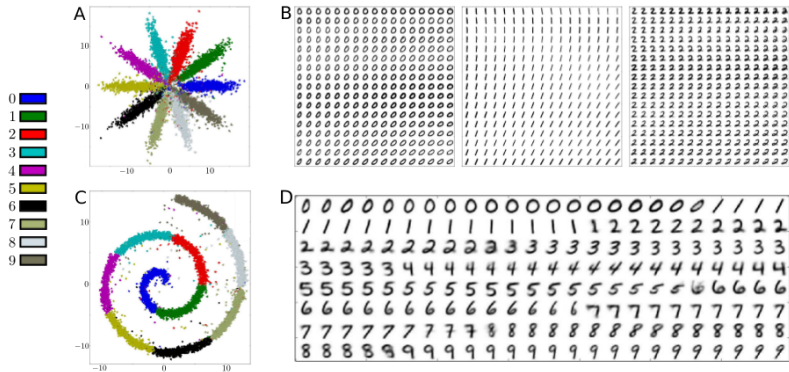


Figure 4: Leveraging label information to better regularize the hidden code. **Top Row:** Training the coding space to match a mixture of 10 2-D Gaussians: (a) Coding space z of the *hold-out* images. (b) The manifold of the first 3 mixture components: each panel includes images generated by uniformly sampling the Gaussian percentiles along the axes of the corresponding mixture component. **Bottom Row:** Same but for a swiss roll distribution (see text). Note that labels are mapped in a numeric order (i.e., the first 10% of swiss roll is assigned to digit 0 and so on): (c) Coding space z of the *hold-out* images. (d) Samples generated by walking along the main swiss roll axis.

Adversarial Variational Bayes

Adversarial Variational Bayes

Adversarial Variational Bayes replaces inference model $Q(Z | X)$ by a black box:

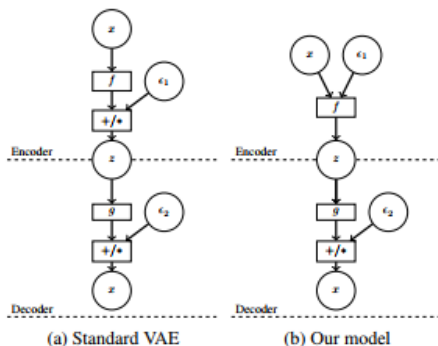


Figure 2. Schematic comparison of a standard VAE and a VAE with black-box inference model, where ϵ_1 and ϵ_2 denote samples from some noise distribution. While more complicated inference models for Variational Autoencoders are possible, they are usually not as flexible as our black-box approach.

Adversarial Variational Bayes

$$\begin{aligned}\mathcal{L}(\theta, \phi) = & \mathbb{E}_x \left[-\text{KL}(q_\phi(z | x) \| p(z)) + \mathbb{E}_{q_\phi(z|x)} \log p_\theta(x | z) \right] = \\ & \mathbb{E}_x \mathbb{E}_{q_\phi(z|x)} [\log p(z) - \log q_\phi(z | x) + \log p_\theta(x | z)]\end{aligned}$$

Solution for the optimization problem:

$$\begin{aligned}\max_T \mathbb{E}_x \mathbb{E}_{q_\phi(z|x)} \log \sigma(T(x, z)) + \mathbb{E}_x \mathbb{E}_{p(z)} \log(1 - \sigma(T(x, z))) \\ T^*(x, z) = \log q_\phi(z | x) - \log p(z)\end{aligned}$$

$$\boxed{\mathcal{L}(\theta, \phi) = \mathbb{E}_x [-T^*(x, z) + \log p_\theta(x | z)]}$$

Adversarial Variational Bayes

AVB = VAE objective +
black box inference +
adversarial estimation of KL;

Summary

Varitional AutoEncoder:

- AutoEncoder that drives code space to match predefined distribution;
- often inferior to GAN;
- stable and relatively fast training;
- various adversarial modifications.

References I

- Doersch, C., 2016. Tutorial on variational autoencoders. arXiv preprint arXiv:1606.05908.
- Kingma, D.P. and Welling, M., 2013. Auto-encoding variational bayes. arXiv preprint arXiv:1312.6114.
- Dumoulin, V., Belghazi, I., Poole, B., Lamb, A., Arjovsky, M., Mastropietro, O. and Courville, A., 2016. Adversarially learned inference. arXiv preprint arXiv:1606.00704.

References II

- Donahue, J., Krähenbühl, P. and Darrell, T., 2016. Adversarial feature learning. arXiv preprint arXiv:1605.09782.
- Makhzani, A., Shlens, J., Jaitly, N., Goodfellow, I. and Frey, B., 2015. Adversarial autoencoders. arXiv preprint arXiv:1511.05644.
- Mescheder, L., Nowozin, S. and Geiger, A., 2017. Adversarial variational bayes: Unifying variational autoencoders and generative adversarial networks. arXiv preprint arXiv:1701.04722.