

Practical tricks

Machine Learning and Data Mining

Maxim Borisyak

National Research University Higher School of Economics (HSE)

Outline

Here, we consider practical problems that are not quite aligned with theory:

- ❑ imbalanced datasets;
- ❑ differences in training and application domains;
- ❑ one-class classification.

Imbalanced datasets

Imbalanced datasets

Settings:

- ❖ classification problem: \mathcal{C}^+ against \mathcal{C}^- ;
- ❖ often in practice $P(\mathcal{C}^+) \ll P(\mathcal{C}^-)$.

This poses several problems:

- ❖ mini-batch learning procedures degradate;
 - ❖ extremely slow learning;
- ❖ imprecise results.

Degradation of mini-batch learning

Probability of a example from \mathcal{C}^+ being selected into a mini-batch is low:

- ❑ \Rightarrow increased $\mathbb{D}[\nabla\mathcal{L}]$;
- ❑ \Rightarrow low learning rate;
- ❑ \Rightarrow slow learning.

Changing priors

$$P(\mathcal{C}^+ | X) = \frac{P(X | \mathcal{C}^-)P(\mathcal{C}^-)}{P(X | \mathcal{C}^-)P(\mathcal{C}^-) + P(X | \mathcal{C}^+)P(\mathcal{C}^+)};$$

$$\frac{L^+}{L^-} = \frac{P(\mathcal{C}^+ | X)}{P(\mathcal{C}^- | X)} = \frac{P(X | \mathcal{C}^+)}{P(X | \mathcal{C}^-)} \cdot \frac{P(\mathcal{C}^+)}{P(\mathcal{C}^-)}.$$

Let $\mathcal{D}[h] = \{\{x | h(x) > \tau\} | \tau \in \mathbb{R}\}$ i.e. set of decision surfaces.

$f(x) = P(\mathcal{C}^+ | X)$ - the ideal classifier for \mathcal{C}^+ against \mathcal{C}^- .

$$\mathcal{D}[f] = \mathcal{D}\left[\frac{f}{1-f}\right] = \mathcal{D}\left[\frac{L^+}{L^-}\right] = \mathcal{D}\left[\frac{P(X | \mathcal{C}^+)}{P(X | \mathcal{C}^-)}\right]$$

Changing priors

An ideal classifier:

- ❖ invariant to change of priors w.r.t. set of decision surfaces;
- ❖ change of priors might stabilize and speed up learning;
- ❖ **only true for a really good classifiers!**

Not ideal classifier under change of priors:

- ❖ low capacity classifiers might change surfaces significantly;
- ❖ dramatic changes of priors may render classifier useless.

Sampling

- ❖ slow convergence is a result high gradient variance;
- ❖ which is the result of unstable composition of mini-batches.

Solutions:

- ❖ stratified batches (forcing class ratio into batches);
- ❖ change in sampling distribution (importance sampling).

Importance sampling

$$\mathcal{L} = \mathbb{E}_{x,y \sim P_{x,y}} l(f(x), y) =$$

$$\int_{x,y} P(x, y) l(f(x), y) dx dy =$$

$$\int_{x,y} P'(x, y) \frac{P(x, y)}{P'(x, y)} l(f(x), y) dx dy =$$

$$\mathbb{E}_{x,y \sim P'_{x,y}} w(x, y) l(f(x), y).$$

$$w(x, y) = \frac{P(x,y)}{P'(x,y)} - \text{weights.}$$

Importance sampling

Resampling trick allows to:

- ❖ stabilize increase frequency of rare but important samples;
- ❖ increase speed of convergence.

Importance can be predetermined e.g.:

- ❖ uniform sampling across classes;
- ❖ increased sampling probability of hard examples.

Weights can be computed on the fly e.g. adaptive sampling methods.

Reweighting

Reweightings

Settings:

- ❖ training set X with distribution P ;
- ❖ target set X' with distribution P' but with absent targets;
- ❖ $P(x) \neq P'(x)$, but
- ❖ $\text{supp } P = \text{supp } P'$.

Examples:

- ❖ training on results of computer simulations.

Reweighting

- ❖ train a classifier $r(x)$ on X' against X :

$$w(x) = \frac{r(x)}{1 - r(x)} = \frac{P'(x)}{P'(x) + P(x)} \cdot \frac{P'(x) + P(x)}{P(x)} = \frac{P'(x)}{P(x)}$$

- ❖ use output as weights (similar to importance sampling):

$$\mathcal{L}_{\text{target}} = \mathbb{E}_{x,y \sim P'} l(f(x), y) = \mathbb{E}_{x,y \sim P} w(x) l(f(x), y)$$

BDT reweighting

Boosting training scheme allows for an especially efficient reweighting algorithm:

- ❖ $w^0(x) = 1$
- ❖ repeat until new classifier yield random guess performance:
 - ❖ train new classifier f^t on X' against X with weights $w^t(x)$;
 - ❖ $w^{t+1}(x) = w^t(x) \frac{f^t(x)}{1-f^t(x)}$.

Semi-supervised learning

Settings

Semi-supervised learning targets cases with a large amount of unlabeled data:

- ❖ $\mathcal{D}_{\text{supervised}} = \{(x_i, y_i)\}_{i=1}^N$;
- ❖ $\mathcal{D}_{\text{unsupervised}} = \{x_i\}_{i=1}^M$;
- ❖ $|\mathcal{D}_{\text{unsupervised}}| \gg |\mathcal{D}_{\text{supervised}}|$;
- ❖ distributions of X are equal in both datasets.
- ❖ also can be used for unbalanced datasets.

Semi-supervised learning

Common techniques:

- ❖ train dimensionality reduction method on $\mathcal{D}_{\text{unsupervised}}$;
- ❖ apply dimensionality reduction to $\mathcal{D}_{\text{supervised}}$;
- ❖ solve supervised problem reduced domain.

Examples:


- ❖ feature selection + classifier;
- ❖ PCA + classifier.

Semi-supervised learning

Dimensionality reduction and supervised methods are trained independently:

- ❖ conflict of objectives:
 - ❖ reducing dimensionality without losses \neq easier supervised problem;
 - ❖ information lost in compression might be important for supervised task.

Conflict of objectives



Semi-supervised deep learning

In Deep Learning both objectives (dimensionality reduction and supervised task) can be trained simultaneously, e.g.:

- ❖ encoder $z = e(x)$;

- ❖ decoder $x' = d(z)$;

- ❖ classifier $f(z)$

$$\mathcal{L} = \mathbb{E}_{X,Y \sim \text{supervised}} l_1(f(e(x)), y) + \lambda \mathbb{E}_{X \sim \text{unsupervised}} l_2(x, d(e(x)))$$

One-class classification

Settings

- ❖ training dataset consist only from one class \mathcal{C}^+ ;
- ❖ target dataset might contain additional classes.

Examples:

- ❖ anomaly detection;
- ❖ outlier detection;
- ❖ novelty detection.

One-class classification

Essentially, generative problem:


- ❖ decision function: $P(X|\mathcal{C}^+) > \tau$:
- ❖ however, objective is a classification one w.r.t unknown class;
- ❖ thus, ill-defined problem.

One-class SVM

Minimizes volume contained by class:

$$\begin{aligned} \min \quad & \frac{1}{2} \|w\|^2 \quad + \quad \frac{1}{\nu n} \sum_i [\xi_i - \rho] \\ \text{subject to} \quad & \\ w\phi(x_i) \geq \rho - \xi_i \quad & : \quad \forall i \\ \xi_i \geq 0 \quad & : \quad \forall i \end{aligned}$$

One-class SVM



Source: sklearn

Dimensionality reduction

The following heuristic might help:

- ❖ train an Auto-Encoder e, d on positive class;
- ❖ compute distribution of reconstruction errors $P[(x - d(e(x)))^2]$;
- ❖ use this distribution as score for one-class classification:

$$P[(x - d(e(x)))^2] > \tau$$

- ❖ Auto-Encoder should be heavily restricted;
- ❖ better to use denoising AE:

$$\sum_i [x_i - d(e(x_i + \varepsilon))]^2 \rightarrow \min$$

Restricted networks

Some use a network:

- ❖ $f(x)$ trained to replicate $y(x) = 1$;
- ❖ f is heavily restricted:
 - ❖ bottleneck does not allow to learn $y(x) = 1$ for all x .

Examples:

- ❖ Radial Basis Networks:

$$f(x) = \sum_i w_i \exp(-\|x - c_i\|^2)$$

One against everything

One against everything = semi-supervised + one-class:

- ❖ large unlabeled dataset \mathcal{D} ;
- ❖ small positive dataset \mathcal{C}^+ ;
- ❖ train positive class against everything:

$$f(x) = \frac{P(X | \mathcal{C}^+)}{P(X | \mathcal{C}^+) + P(X | \mathcal{D})} \sim P(X | \mathcal{C}^+)$$

One against everything

Examples:

- ❖ it is easy to sample large amounts of text (e.g. tweets);
- ❖ sampling abnormal text might be problematic.

Summary

Summary

Imbalanced datasets:

- ❖ be careful with changing priors;
- ❖ resampling.

Importance sampling:

- ❖ may improve convergence and stability;
- ❖ importance sampling optimization improves convergence rate.

Reweighting:

- ❖ different training and target distributions of X .
- ❖ a special case of domain adaptation.

One-class classification:

- ❖ a very strange field;
- ❖ ill-defined generative problem.